

1. _____ is a framework that allows distributed processing of large datasets across clusters of computers using simple programming models.
 - A) Hadoop
 - B) Big data
 - C) Distributed systems
 - D) Networking
2. When a file is split into blocks, the default size of each file block is _____.
 - A) 12 GB
 - B) 156 MB
 - C) 128 MB
 - D) 30 MB
3. Creating copies of a block in multiple DataNodes is called _____.
 - A) Copying
 - B) Block Replication
 - C) Data Block Split
 - D) Data Log
4. A _____ is a map-side join where every single record is paired up from another dataset.
 - A) Reduce side join
 - B) Composite join
 - C) Cartesian product
 - D) Replicated join
5. Which of the following statements is NOT a characteristic of MapReduce?
 - A) Works well on write once read many (WORM) data
 - B) Operations are performed by multiple processors
 - C) Handles small-scale data
 - D) Leverages commodity hardware and storage
6. When a table is created in Hive, data is stored as text with _____ comma-separated fields per line.
 - A) Two
 - B) Four
 - C) Three
 - D) Five
7. Which of the following is NOT a primitive data type in Hive?
 - A) ARRAY
 - B) BOOLEAN
 - C) FLOAT
 - D) DOUBLE
8. _____ is a command-line interface application for transferring data between relational databases and Hadoop.
 - A) MapReduce
 - B) SQOOP
 - C) Netezza
 - D) MySQL
9. What action does the 'Impala-shell ?q ?select * from simple?' command perform?
 - A) Runs direct queries from shell using the ?q option
 - B) Checks all options of Impala using the help option
 - C) Issues a use database on startup using the ?d option
 - D) Runs direct queries from shell using the ?d option
10. The character "_____" is used to execute Beeline commands.
 - A) #
 - B) @
 - C) *

D) !

11. Which of the following features does Hive NOT include?
 - A) Comprises a specialized SQL engine
 - B) Used mostly for batch processing
 - C) Used mainly for interactive queries and data analysis
 - D) Accommodates many concurrent users
12. Which of the following is NOT part of the query execution process in Impala?
 - A) Make optimizations
 - B) Process data using MapReduce or Apache Spark
 - C) Execute query on cluster
 - D) Store the data in HDFS
13. By 2020, the amount of new information created for each human being will be _____.
 - A) 100 MB
 - B) 1.7 MB
 - C) 1 GB
 - D) 50 GB
14. Using multiple machines to execute a job is called _____.
 - A) Multiple server system
 - B) LAN
 - C) WAN
 - D) Distributed systems
15. How long does it take for one machine to process 1 Terabyte of data?
 - A) 1 hour
 - B) 30 minutes
 - C) 45 minutes
 - D) 1 minute