

1. What action does the command line ``$hdfs dfs -ls`` perform?
 - A) Copies file to the user's directory
 - B) Creates a directory in the user's home directory
 - C) Deletes the directory
 - D) Gets a directory listing of user's home directory in HDFS
2. MapReduce was introduced by _____.
 - A) Microsoft
 - B) Google
 - C) Oracle
 - D) IBM
3. Which of the following is NOT a real-time use of MapReduce?
 - A) Algorithms
 - B) Gaussian Analysis
 - C) Data Transfer
 - D) Enterprise Analytics
4. A _____ is a map-side join where every single record is paired up from another dataset.
 - A) Reduce side join
 - B) Composite join
 - C) Cartesian product
 - D) Replicated join
5. Which of the following statements is NOT a characteristic of MapReduce?
 - A) 2 and 3
 - B) 1 and 4
 - C) 1, 2, 3, 4
 - D) 2, 3, 4
6. The MapReduce responsibilities of a developer does NOT include _____.
 - A) Setting up the job
 - B) Distributing the job
 - C) Specifying the input location
 - D) Ensuring correct format and location of input
7. _____ is a Hadoop feature that helps cache files needed by applications.
 - A) Distributed Cache
 - B) Distributed Systems
 - C) Distributed Processing
 - D) Distributed Ecosystem
8. _____ is a command-line interface application for transferring data between relational databases and Hadoop.
 - A) MapReduce
 - B) SQOOP
 - C) Netezza
 - D) MySQL
9. What action does the `'Impala-shell ?q ?select * from simple?'` command perform?
 - A) Runs direct queries from shell using the `?q` option
 - B) Checks all options of Impala using the help option
 - C) Issues a use database on startup using the `?d` option
 - D) Runs direct queries from shell using the `?d` option
10. The character "_____" is used to execute Beeline commands.
 - A) #
 - B) @
 - C) *
 - D) !
11. Which of the following features does Hive NOT include?

- A) 1, 2, 3, 4
- B) 2, 3, 4
- C) 1, 2, 3
- D) 1, 3, 4

12. Which of the following is NOT part of the query execution process in Impala?

- A) Make optimizations
- B) Process data using MapReduce or Apache Spark
- C) Execute query on cluster
- D) Store the data in HDFS

13. When a file is split into blocks, the default size of each file block is _____.

- A) 12 GB
- B) 156 MB
- C) 128 MB
- D) 30 MB

14. Creating copies of a block in multiple DataNodes is called _____.

- A) Copying
- B) Block Replication
- C) Data Block Split
- D) Data Log

15. Which of the following access mechanisms does HDFS not provide?

- A) Java
- B) Python
- C) Web GUI
- D) C