

Here are your MCQs in A, B, C, D format (without answers):

1. When a file is split into blocks, the default size of each file block is _____.
 - A) 12 GB
 - B) 156 MB
 - C) 128 MB
 - D) 30 MB
2. By default, all data in the data warehouse directory structure is stored in _____.
 - A) /user/impala/warehouse
 - B) /hive/user/warehouse
 - C) /user/hive/warehouse
 - D) /user/impala/rdbms
3. When a table is created in Hive, data is stored as text with _____ comma-separated fields per line.
 - A) Two
 - B) Four
 - C) Three
 - D) Five
4. Which of the following is NOT a primitive data type in Hive?
 - A) ARRAY
 - B) BOOLEAN
 - C) FLOAT
 - D) DOUBLE
5. _____ is a command-line interface application for transferring data between relational databases and Hadoop.
 - A) MapReduce
 - B) SQOOP
 - C) Netezza
 - D) MySQL
6. What action does the 'Impala-shell ?q ?select * from simple?' command perform?
 - A) Runs direct queries from shell using the ?q option
 - B) Checks all options of Impala using the help option
 - C) Issues a use database on startup using the ?d option
 - D) Runs direct queries from shell using the ?d option
7. The character "_____" is used to execute Beeline commands.
 - A) #
 - B) @
 - C) *
 - D) !
8. While validating data, missing data is represented as _____.
 - A) ZERO
 - B) ERROR
 - C) NULL
 - D) MISSING
9. Which of the following features does Hive NOT include?
 - A) 1, 2, 3, 4
 - B) 2, 3, 4
 - C) 1, 2, 3
 - D) 1, 3, 4
10. Which of the following is NOT part of the query execution process in Impala?
 - A) Make optimizations
 - B) Process data using MapReduce or Apache Spark

- C) Execute query on cluster
- D) Store the data in HDFS

11. Match the following:

- A) A1-B3, A2-B2, A3-B1
- B) A1-B2, A2-B3, A3-B1
- C) A1-B1, A2-B3, A3-B2
- D) A1-B3, A2-B1, A3-B2

12. By 2020, the amount of new information created for each human being will be _____.

- A) 100 MB
- B) 1.7 MB
- C) 1 GB
- D) 50 GB

13. Using multiple machines to execute a job is called _____.

- A) Multiple server system
- B) LAN
- C) WAN
- D) Distributed systems

14. How long does it take for one machine to process 1 Terabyte of data?

- A) 1 hour
- B) 30 minutes
- C) 45 minutes
- D) 1 minute

15. _____ is a framework that allows distributed processing of large datasets across clusters of computers using simple programming models.

- A) Hadoop
- B) Big data
- C) Distributed systems
- D) Networking